

## MF810: Environment setup and fundamentals

---

## Desirable properties of a program or code

Some important features of a program and its source code:

**Efficiency** The most obvious: Do what we want without taking too long.

**Correctness** Do what we want in all circumstances, including edge cases.

Note that in unexpected cases it is often better to abort than to just continue.

**Robustness** Handle edge cases, unexpected inputs, or errors gracefully.

**Maintainability** Adding features, changing features, or fixing bugs should be easy.

**Readability** Somebody else (or future you) should be able to understand the code.

Not only about **what** the code does, but also **why**.

Helps maintainability.

**Portability** The program should run correctly on other computers than the developer's.

- Docker and Linux basics
- Python Docker setup
- Intro to Pytorch
- Intro to neural networks
- Neural neural networks in Pytorch

## Portability: Docker and Linux

- When moving code from one computer to another, it can fail to run. Some reasons include
  - Dependencies are not installed. *package X exist on another*
  - Package versions differ.
  - Files are not in the right location.
  - Etc.
- These issues are solved by **containers**. (They also provide some safety/security benefits if used correctly.) *↑ solve problem*
- **Docker** is a system for running containers.
- Docker is built on features of the **Linux** kernel. To be proficient Docker users, we must build basic understanding of Linux topics.

# What is Linux?

## Linux

- **Linux** typically refers to an operative system, like Windows or macOS.
- Strictly speaking, Linux is the **Linux kernel**.
- A kernel is the piece of software that coordinates everything in the background and communicates with the hardware. As far as software is concerned, the kernel is the computer!
- Examples include: the *Windows NT Kernel* on Windows; *XNU* on Apple devices; the *Linux kernel* on Android, etc.
- The Linux kernel provides features to run kernels inside the kernel itself,<sup>1</sup> which is what makes Docker and containerization possible.

---

<sup>1</sup>This is technically not what happens, but it is *similar* to doing so.

Let us look at basics of how a Linux environment works.

Windows and macOS are in many ways very similar, but details differ.

We will discuss

- Users, groups, and permissions.
- Filesystems and file permissions.
- Environment variables.
- The “command line.”

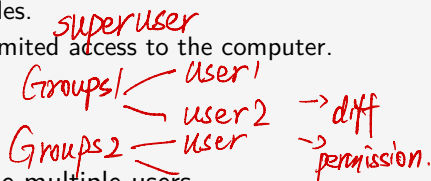
# Users and groups

- When setting up a new computer, you must enter a (user)name. In this process, the OS sets up a **user** (account) for you.
- Systems typically have one super user<sup>2</sup> and one or more regular users.
- The purpose of multiple users is to restrict users from doing certain things:
  - One users should not be able to read other users' files.
  - Only the super user (admin) should be granted unlimited access to the computer.
  - Etc.

- Different users have different **permissions**.

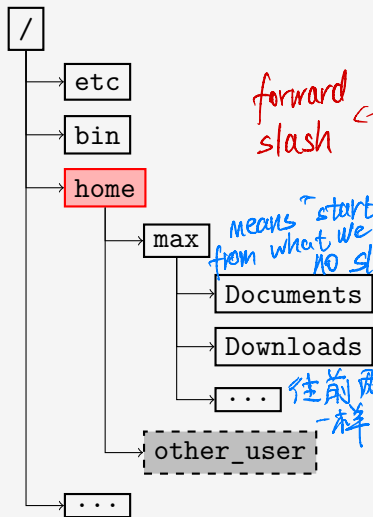
- Users can be members of **groups**, which may include multiple users.

Groups are another tool to grant permissions. For instance, users who are members of the printer group are allowed to use the printers.



<sup>2</sup>**root** user on Linux/Unix/macOS or *Administrator* on Windows.

# Filesystems



- The **filesystem** is a sequence of **directories** (folders) starting at the **root** (`/` on Linux/Unix/macOS; like `C:/` on Windows).

- A file location is given by a **path**:

`/home/max/Documents/MF810/lecture1.pdf`

This is called an **absolute path** because it references the file from the root.

- A **relative path** is relative to the current directory. From

`/home/max` this relative path is the same file:

`Documents/MF810/lecture1.pdf`

- Special path components: `./` means same/current directory;

`../` means **parent directory** ("up/back").

- From `/home/max/Downloads`, these are still the same files:

`../Documents/MF810/lecture1.pdf` and

`../../Documents/MF810/./lecture1.pdf`

- Relative paths should be used within a project!



## Files and file permissions

- A directory may contain a number of **files**.
- There are three main types of files: regular files, symlinks (references/shortcuts), and directories.
- Each file has an owner, a group owner, and two sets of permissions: owner permissions and group permissions.
- A file listing may look like

-rw-r--r-- 1 max users 206K Jan 13 14:23 lecture1.pdf  
type (-,l,d) permissions owner group size date filename

*owner*

rw- indicates that the owner (max) may read (r) and write (w) to the file.

*other*

r-- indicates that all members of the group users may read (r) but not write (w) to the file.

*run(x)*

- Permissions are central to the system and a great protection against accidental mistakes.

# Docker images

- Docker images and containers are like classes and objects in object-oriented programming: containers are specific instances of images.
- Images can be created in two ways:
  - By downloading an existing image, e.g.  
`docker pull ubuntu` or `docker pull python:slim`.
  - By extending an existing image in a `Dockerfile` and building it:  
`docker build -t image_name path_to_Dockerfile_dir`  
We will use the name `my_python_im` and the path will be `.` if run from the same directory as the `Dockerfile`.
- Containers are created from an image using either `docker run` or `docker create`.
- Making changes to a container does not change the image it was created from!

## Containers: `docker run` example

- Once set up, we run any Python script by entering

```
docker run --rm -it my_python_im -v "$PWD":/workdir python scriptname.py
```

- `docker run` creates and runs a container from an **image**.
- `--rm` deletes the container after termination.
- `-it` specifies interactive mode (as opposed to running in the background).
- `my_python_image` is the name we have given our Python image.
  - `-v "$PWD":/workdir` makes the current directory (`$PWD`) available in the container at the `/workdir` path.
  - `python` specifies that the container should execute the Python executable.
  - `scriptname.py` is the script file we want python to run.
- This is long, but we can make a shortcut as e.g.  
`dockpy scriptname.py`
- Like this, your code can be sent to a friend or a **cloud computer** and run exactly like it does on your computer!

# Deep learning libraries and neural networks

---

- What are famous libraries like Tensorflow and Pytorch?
- Short answer:  
They are numerical libraries with extra support for differentiating special functions.
- Consider an abstract problem of minimizing a function  $L$  as a function of  $\theta$ :

$$\min_{\theta \in \mathbb{R}^P} L(\theta).$$

- If  $L$  is smooth, the first order condition at the minimum states \* --> 表示 optimal.

$$\nabla_{\theta} L(\theta) \Big|_{\theta=\theta^*} = \nabla_{\theta} L(\theta^*) = 0.$$

- Gradients are intimately connected to optimization, and that is the interest in numerical libraries with differentiation support.


# Gradient descent

- Recall from calculus, that the gradient  $\nabla_{\theta} L$  is a **vector** pointing in the direction in which the function  $L$  is most increasing (steepest).
- Similarly,  $-\nabla_{\theta} L = \nabla_{\theta}(-L)$  is the direction in which  $-L$  is most increasing, i.e.,  $-\nabla_{\theta} L$  is the direction that  $L$  is most decreasing.
- If we from any point  $\theta$  move a small distance  $\epsilon \nabla_{\theta} L(\theta)$  along the gradient to  $\theta'$ , we should have  $L(\theta') \leq L(\theta)$ .
- We have discovered **gradient descent**:

$$\left[ \theta_{n+1} = \theta_n - \epsilon \nabla_{\theta} L(\theta_n). \right]$$

*at point* *minus* *move a small distance.*  
 $\epsilon \rightarrow$  learning rate  
 $\epsilon$

- Because our deep learning library can compute gradients, this is trivial to implement!
- The parameter  $\epsilon$  is called the learning rate.



## Our first example: linear regression

- Consider the problem of least squares linear regression:

$x$  is vector  
 $x_i, i \in \{1, \dots, N\}$   
 $\in \mathbb{R}^p$

where

$$\min_{\theta \in \mathbb{R}^p} \sum_{i=1}^N \ell(y_i, f(x_i; \theta))$$

$\ell(y, y') = |y - y'|^2$  and  $f(x; \theta) = \theta \cdot x = \sum_{j=1}^p \theta_j x_j$

*inner product*

- This is called linear regression because  $f$  is linear in  $\theta$ , not because it is linear in  $x$ !
- Define

$$L(\theta) = \sum_{i=1}^N \ell(y_i, f(x_i; \theta))$$

same as  $\min_{\theta} L(\theta)$   
so take gradient

so that

$$\nabla_{\theta} L(\theta) = \sum_{i=1}^N \partial_{y'} \ell(y_i, f(x_i; \theta)) \nabla_{\theta} f(x_i; \theta) = \sum_{i=1}^N 2(y_i - x_i \cdot \theta) x_i.$$

- These expressions are much simpler on matrix form.

# Linear regression gradient descent

- With

$$\nabla_{\theta} L(\theta) = \sum_{i=1}^N 2(y_i - x_i \cdot \theta)x_i$$

gradient descent is equivalent to iterating

$$\theta_{n+1} = \theta_n - \epsilon \sum_{i=1}^N 2(y_i - x_i \cdot \theta)x_i.$$

*start with* (pointing to  $\theta_n$ )  
*times with* (pointing to the sum term)

- For linear regression, this method converges to the minimizer  $\theta^*$ .<sup>3</sup>

---

<sup>3</sup>If  $\epsilon$  is decreasing at an appropriate rate.



- The function  $f(x; \theta) = \theta \cdot x = \theta x$  maps  $x \in \mathbb{R}^P$  to  $\mathbb{R}$ .
- Generalize this to  $\mathbb{R}^P$  to  $\mathbb{R}^k$  by taking  $\theta \in \mathbb{R}^{P \times k}$ .

$$\mathbb{R}^P \rightarrow \mathbb{R}^k \quad x \mapsto \theta x.$$

- It is useful to apply another function  $\sigma$  to the output.
- Often  $\sigma : \mathbb{R} \rightarrow \mathbb{R}$ , in which case this is done element-by-element, and we write  $\sigma$  for such **broadcasting**:

$$\sigma.(x) = (\sigma(x_1), \dots, \sigma(x_k)), \quad x \in \mathbb{R}^k.$$

We write just  $\sigma(x)$  to mean  $\sigma.(x)$  in this case.

- The transformation  $x \mapsto \theta x$  is called a **linear layer**.
- The function  $f(x; \theta) = \sigma(\theta x)$  is called just a **layer** or a **dense layer**.

- The function  $x \mapsto \sigma(\theta x)$  is an example of a trivial (**artificial**) **neural network** (**ANN** or **NN**).

- Let  $\theta = (w_1, w_2)$  and define  $h_1(x; w_1) = \sigma_1(w_1 x)$  and  $h_2(x; w_2) = \sigma_2(w_2 x)$ .  
The composition (assuming  $w_1 \in \mathbb{R}^{k_0 \times k_1}$  and  $w_2 \in \mathbb{R}^{k_1 \times k_2}$ )

$$(h_2 \circ h_1)(x) = h_2(h_1(x)) = \sigma_2(w_2 \sigma_1(w_1 x))$$

is a **two-layer neural network**.

- The function  $h_1$  is called a **hidden layer** and the function  $h_2$  is called the **output layer**.
- The parameters  $w_1$  and  $w_2$  are called **weights**.
- The functions  $\sigma_1$  and  $\sigma_2$  are called **activation functions**.

## Neuron bias

- For reasons we discuss later, it is common to work with affine transformations instead of linear.
- Let  $\theta = ((w_1, b_1), \dots, (w_d, b_d))$  and define *each adding b.*

$$h_i(x) = \sigma_i(\underline{b_i} + w_i x).$$

- The term  $b_i$  is called a **neuron bias**, often **bias** for short.
- The composition  $h_d \circ h_{d-1} \circ \dots \circ h_2 \circ h_1$  is a **neural network** of **depth  $d$** .
- Because  $h_d(x)$  is often a linear transformation to a scalar ( $\mathbb{R}^{k_{d-1}} \rightarrow \mathbb{R}^{k_d} = \mathbb{R}$ ), it is customary to speak of the number of **hidden layers**, which here is  $d - 1$ .
- Like before, the weights must satisfy  $w_i \in \mathbb{R}^{k_{i-1} \times k_i}$  and  $b \in \mathbb{R}^{k_i}$ .

The value  $k_i$  is called the **width** of layer  $i$ .

- The number of parameters in this network is

$$P = \sum_{i=1}^d (k_{i-1} + 1) k_i.$$

*depth \* width ?*

- This type of neural network is called a **multilinear perceptron (MLP)**.
- An MLP is a **feedforward** neural network in which all neurons in two adjacent layers are connected to each other (called **fully connected** or **dense**).
- There are many other examples of layers. Some examples include **Convolutional layers** Feedforward but not fully connected.



Commonly used to find spatial patterns, like in images.

**Recurrent layers** Not feedforward.

Used for sequential data like natural language processing (NLP) and time series.

**Dropout layers** Randomly omits connections: feedforward but not dense.

Used for regularization.

- Such common structures are included in the libraries and can be used directly.
- There are no restrictions on general layers; custom layers are often created.